



**School of Computer Science
Faculty of Science
COMP-2560: System Programming (Fall 2024)**

Lab#	Date	Title	Due Date	Grade Release Date
Lab07	Week 07	Devices as Files	One-Week Lab Oct. 29 Tuesday Midnight EDT	Nov. 4

The main objective of this lab will be to learn how to UNIX handles devices as files. For the sake of this lab, whenever we refer to a file that is a device, we say *special* file. Otherwise, it is a *normal* file like an image file or program file.

Step 1. UNIX Devices as Special Files

UNIX's kernel at its high-level of File System (File Manager) defines *file* as an abstract entity to read from or write to. File could be part of a hard disk drive that stores your C program `main.c`, another part of the hard disk drive that stores your program executable `main`, or the whole hard disk drive. UNIX's kernel at its File System assumes all other hardware devices are also files. The differences between the devices are handled by another part of the kernel, Device Manager. So, like your files in your home directory, devices are also files at `/dev` directory:

```
hfani@alpha:~$ ls /dev
autofs      fb0          kvm           ptmx         shm          tty14       tty26
block       fd           log           ptp0         snapshot    tty15       tty27
bsg         full         loop-control  ptp1         snd         tty16       tty28
btrfs-control fuse         mapper        pts          stderr      tty17       tty29
bus         hidraw0      megaraid_sas_ioctl_node random        stdin       tty18       tty3
char        hidraw1      mem           rfkill       stdout      tty19       tty30
console     hpet         mqueue       rtc          tty         tty2        tty31
core        hugepages   net           rtc0         tty0        tty20       tty32
cpu         hwrng       null          sda          tty1        tty21       tty33
cpu_dma_latency initctl      nvram         sda1         tty10       tty22       tty34
cuse        input       port          sda2         tty11       tty23       tty35
disk        ipmi0       ppp           sda5         tty12       tty24       tty36
dri         kmsg        psaux         sg0          tty13       tty25       tty37
```

As seen, we have `cpu` (processor) and `mem` (memory), which are the two essential parts of any computer system. School's computer system is a multiprocessor system; that is, it has multiple cpus. So, you can go inside the `/dev/cpu` and see how many cpus exist:

```
hfani@alpha:~$ ls /dev/cpu
0 1 10 11 12 13 14 15 2 3 4 5 6 7 8 9
```

Let's go into the first cpu:

```
hfani@alpha:~$ ls /dev/cpu/0
cpuid msr
```

Amazingly, devices are managed like files and directories that are directly connected to the devices. You may be tempted to open them as normal files using `vi`:

```
hfani@alpha:~$ vi /dev/cpu/0/cpuid
"/dev/cpu/0/cpuid" is not a file
```

You see that `vi` is not able to read from `cpu/0/cpuid` like a normal file. Let's try it for memory device:

```
hfani@alpha:~$ vi /dev/mem
"/dev/mem" is not a file
```

Is this because devices are not files and what we say is incorrect, or `vi` program is not able to read from these *special* files? We'll answer this shortly.

What are other files? What is the file for mouse, keyboard, monitor, sound, mic, hard disk drives, etc?

Important parts of a computer system are storage devices. Depending on the type of storage, the file name for the device is either start with `hd` or `sd` followed by a letter that distinguishes multiple hard disk storages. For instance, `sda` is the first hard disk, `sdb` is the second, and so on. Also, each hard disk storage can have different partitions, which are identified by a number. For instance, `sda1` is the first partition of the first hard disk storage. As you can see, the School's system has one hard disk storage with 3 partitions.

```
hfani@alpha:~$ ls /dev/sd*
/dev/sda /dev/sda1 /dev/sda2 /dev/sda5
```

Who select these names for devices? Are there any standards? Can I rename the file of a device? For instance, can I change the `/dev/mem` to `/dev/memory`?

Step 2. Standard File Descriptors

There are three files that *look like* they are devices, but they are not. They are file descriptors (numbers) that are connected to either a normal file or a special file (device):

```
hfani@alpha:~$ ls /dev/fd
0 1 2 3
hfani@alpha:~$ ls /dev/std*
/dev/stderr /dev/stdin /dev/stdout
```

Traditionally, `/dev/fd/0` or `/dev/stdin` was connected to keyboard device, meaning that reading from keyboard file (device) was the same as reading from `/dev/fd/0` or `/dev/stdin` file descriptors. `/dev/fd/1` or `/dev/stdout` was connected to printing device or monitor, meaning that writing to printer file (device) was the same as writing to `/dev/fd/1` or `/dev/stdout` file descriptors; likewise, for `/dev/fd/2` and `/dev/stderr` to print out errors. Why not use special files' names directly for keyboard or printer instead of working with these numbers? Very good question!

Working with special files (device) requires somebody to be next to the School's computer system. For instance, if you want to read from the keyboard from its assigned special file:

```
/dev/input/by-path/pci-0000:00:13.0-usb-0:3:1.1-event-kbd
```

somebody should be in the School's server room and press any keys in the keyboard device attached to the School's computer system. This is not convenient for us when we connect remotely to a computer system. Therefore, by default, UNIX-based/like operating system developers came up with the idea of standardizing the input and output for programs regardless of where or what the actual device is. So, they fixed these three file descriptors that can be connected to any special file (device) or normal file. For the program, it does not matter. The program reads from `0` and writes to `1`. If error, the program writes to `2`.

You may connect `0` and `1` and `2`, all to a normal file (how?). Then your program reads input from the file and writes output and errors, if any, to the same file.

But how do we run our program, like the one in Lab03, at School's computer system that reads input and writes outputs? This is handled by either kernel or shell, who attach standard file descriptors to a virtual terminal device called `tty` (TeleTYpewriter) with a unique number for a user connected to the computer system. This virtual device is the window with the black background that shows you the shell command prompt:

```
hfani@alpha:~$ █
```

You can see your virtual terminal device number by using the `tty` command:

```
hfani@alpha:~$ tty
/dev/pts/9
```



This is the device your program reads from or writes to, not the keyboard or monitor at the School's server room or at your home!

When your program, like in Lab03, asks an input value, it reads from `0` which is connected to your `tty`. So, when you type numbers or characters by your keyboard attached to your computer at home, it goes to your `tty`, and then it forwards them to your program. Can you explain what happens when your program writes outputs and you see it in your monitor?

Step 3. Read/Write to Devices as Files via Shell

`vi` is a program to open normal files that are stored in storage devices. But it cannot open special files (devices) and simply raise an error. However, UNIX, which claims that devices are files, must provide the same system calls to work with devices the same way as files. In the next step, we'll see UNIX fulfills its promise. Not only that, now we see that its shell also provides a built-in command `cat` (concatenate) to do the operations on devices and files in the same way.

To read from a normal file or a special file (device), you can simply use `cat` followed by the name of the file/device:

```
hfani@alpha:~$ cat test.c
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="test.txt";
    int fd = creat(filename, 0);
    if(fd==-1){
        printf("error happend!");
    }
}
hfani@alpha:~$ cat /dev/mem
cat: /dev/mem: Permission denied
hfani@alpha:~$ cat /dev/cpu/0/cpuid
cat: /dev/cpu/0/cpuid: Permission denied
hfani@alpha:~$ cat /dev/sdal
cat: /dev/sdal: Permission denied
hfani@alpha:~$
```

As seen, we do not have enough permission to read from some devices. But we have permission to access our assigned virtual terminal device `tty`. Let's find our unique `tty` again:

```
hfani@alpha:~$ tty
/dev/pts/9
```

Then, read from it:

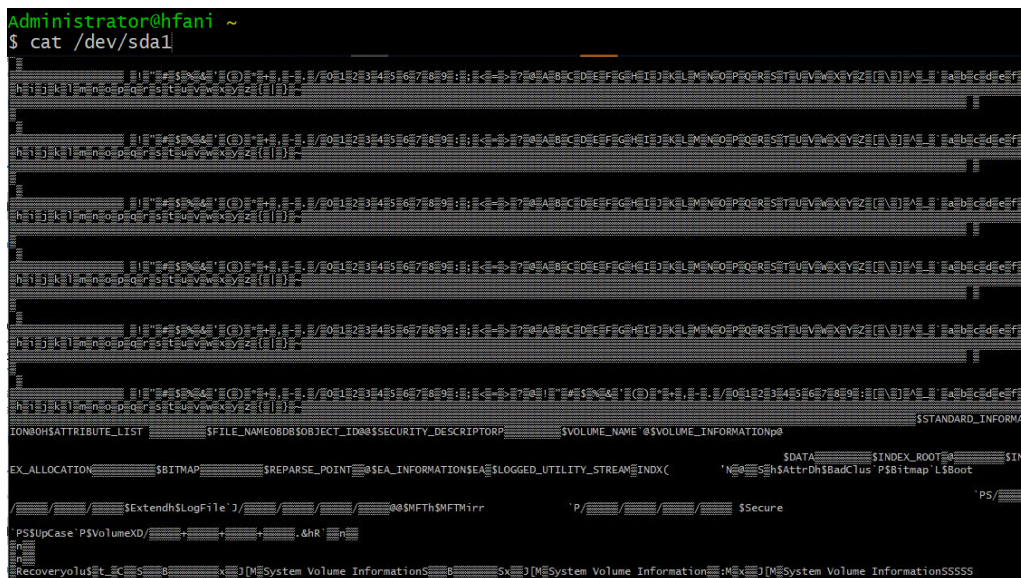
```
hfani@alpha:~$ cat /dev/pts/9
```

As seen, it is empty initially. You can put some data to it by typing and pressing enter:

```
hfani@alpha:~$ cat /dev/pts/9
Hello my virtual terminal device!
Hello my virtual terminal device!
How is it going?
How is it going?
```

You see that the `cat` reads from the `tty` and shows it to us. You can continue to send more info to the `tty` and `cat` reads them and show them to you. You can exit `cat` by `CTRL+C`.

I used `cat` on my first hard disk storage on *my own computer system*, and here is the output:



Let's write to a device. You can use **cat** with **>** operator that connects **/dev** to a device instead of your **tty**. Then, you are able to input data and it writes them to that device:

As seen, you are not able to write to memory due to lack of enough permission.

```
hfani@alpha:~$ cat > /dev/pts/9
Hey my tty again!
Hey my tty again!
I'm back!
I'm back!
```

I don't dare to write on my hard disk storage by `cat!` But I can write to my sound device, whose file is `/dev/dsp`.

```
Administrator@hfani ~  
$ cat > /dev/dsp  
Hello my sound device. would you please sing this for me: "some say Love, it is a rive ..."
```

1) I record my voice in a file `test.wav`

2) I send the content of this file to the sound device

`cat` can connect two files/devices by reading from one file or device (source) and write it into the second file or device (target). For example,

```
hfani@alpha:~$ cat test.c > /dev/pts/9
#include <fcntl.h>
#include <stdio.h>
void main(void){
    char filename[20]="test.txt";
    int fd = creat(filename, 0);
    if(fd== -1){
        printf("error happend!");
    }
}
```

which reads from the file `test.c` and writes it to your `tty` device. Now, I want to send my `test.wav` to my sound device:

```
Administrator@hfani ~
$ cat /cygdrive/c/test.wav > /dev/dsp
```

Great! Please note that we cannot do this on the School's computer system due to lack of permission. You can try it yourself if you have UNIX-base/like OS or UNIX emulator software (e.g., Cygwin¹ on Windows.)

Step 4. Read/Write to Devices as Files via System Calls

Now that we know how devices and files are treated the same in UNIX, let's open a device and read from or write to it like a normal file in a C program. Due to our lack of enough permission on the School's computer system, we use the virtual terminal device `tty` in our example. First, we need to know the `tty` number that assigned to our terminal:

```
hfani@alpha:~$ tty
/dev/pts/9
```

Now, `/dev/pts/9` is our read and write device. First, we need to open it and have a file descriptor to the device.

The `open` system call is in the file control header (`fcntl.h`):

```
#include <fcntl.h>
#include <unistd.h>
#include <stdio.h>
void main(){
    int fd_in = open("/dev/pts/9", O_RDONLY);
```

Then, we need to read from it using the file descriptor at hand. The `read` system call is in UNIX standard header (`unistd.h`):

```
char buf[20];
ssize_t bytes_read;

while ((bytes_read = read(fd_in, buf, sizeof(buf))) > 0) {
    printf("block read: \n<%s>\n", buf);
}
```

Commented [A1]: TTYNAME(3)

Linux Programmer's

Manual YNAME(3) TT

NAME
ttyname, ttyname_r - return name of a terminal

SYNOPSIS
#include <unistd.h>

char *ttyname(int fd);

int ttyname_r(int fd, char *buf, size_t buflen);

DESCRIPTION
The function `ttyname()` returns a pointer to the null-terminated pathname of the terminal device that is open on the file descriptor `fd`, or NULL on error (for example, if `fd` is not connected to a terminal). The return value may point to static data, possibly overwritten by the next call. The function `ttyname_r()` stores this pathname in the buffer `buf` of length `buflen`.

RETURN VALUE
The function `ttyname()` returns a pointer to a pathname on success. On error, NULL is returned, and `errno` is set appropriately. The function `ttyname_r()` returns 0 on success, and an error number upon error.

```
*****
*****
*****
```

```
moradib@charlie:~/Unix2022/lab06$ tty
/dev/pts/3
```

```
moradib@charlie:~/Unix2022/lab06$ ./getttyname
MY TTY name is: /dev/pts/3
```

```
moradib@charlie:~/Unix2022/lab06$ cat get_ttyname.c
```

```
#include<unistd.h>
#include<stdio.h>

int main(int argc, char *argv[])
{
    char *MyTTYName = ttyname(0); /*
Note: ttyname takes Int fd(STDIN_FILENO) */
    printf("MY TTY name is: %s\n", MyTTYName);
    return 0;
}
```

¹ <https://www.cygwin.com/>



Step 5. Lab Assignment

There are two tasks for your submission:

Our example in step 4 never stops. Add lines of code such that when the user writes **END**, it exits the while loop, safely closes the device, and returns.

- 1) Instead of printing what we read from the device to output (the terminal device itself), write them into a normal file, named **log.txt**.

You are not allowed to use any library routines, no `stdio.h`! Only system calls in `fcntl.h` and `unistd.h`. The sample code for step 4 has been attached in a zip file named **lab07_hfani.zip**.

1.1. Deliverables

You will prepare and submit the program in one single zip file **lab07_uwindid.zip** containing the following items:

- (90%) **lab07_uwindid.zip**
- (70%) **tty.c** => built with no error
 - (10%) **log.txt** => the output of the program
 - (10%) **results.pdf/jpg/png** => the image snapshot of the program run
 - (Optional) **readme.txt**

(10%) Files Naming and Formats

Please follow the naming convention as you lose marks otherwise. Instead of uwindid, use your own account name, e.g., mine is **hfani@uwindsor.ca**, so, **lab07_hfani.zip**